

Clasificación de Revenue (Sí/No compra)

integrante: Luis Lucero Balvin

```
# === Colab: Clasificación de Revenue (Sí/No compra) ===

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, StratifiedKFold,
cross_val_score
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, f1_score, roc_auc_score, precision_recall_curve,
    roc_curve, average_precision_score, classification_report,
    confusion_matrix
)
```

1. Carga robusta del CSV

```
# ----- 1) Cargar CSV -----
csv_name = "online_shoppers_intention.csv"
if not os.path.exists(csv_name):
    from google.colab import files
    print(f"No encontré '{csv_name}'. Selecciona tu CSV para
subirlo...")
    uploaded = files.upload()
    if csv_name not in uploaded:
        csv_name = list(uploaded.keys())[0]
        print(f"Usando archivo subido: {csv_name}")

df = None
for enc in ("utf-8", "latin-1", "cp1252"):
    try:
        df = pd.read_csv(csv_name, encoding=enc)
        break
    except Exception:
        pass
if df is None:
    raise RuntimeError("No pude leer el CSV. Verifica
ruta/codificación.")
```

```

print("Dimensiones:", df.shape)
display(df.head())

Dimensiones: (12330, 18)

{
  "summary": {
    "name": "display(df",
    "rows": 5,
    "fields": [
      {
        "column": "Administrative",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": ""
        }
      },
      {
        "column": "Administrative_Duration",
        "properties": {
          "dtype": "number",
          "std": 0.0,
          "min": 0.0,
          "max": 0.0,
          "num_unique_values": 1,
          "samples": [0.0],
          "semantic_type": ""
        }
      },
      {
        "column": "Informational",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": ""
        }
      },
      {
        "column": "ProductRelated",
        "properties": {
          "dtype": "number",
          "std": 3,
          "min": 1,
          "max": 10,
          "num_unique_values": 3,
          "samples": [1],
          "semantic_type": ""
        }
      },
      {
        "column": "ProductRelated_Duration",
        "properties": {
          "dtype": "number",
          "std": 274.5386534209701,
          "min": 0.0,
          "max": 627.5,
          "num_unique_values": 4,
          "samples": [64.0],
          "semantic_type": ""
        }
      },
      {
        "column": "BounceRates",
        "properties": {
          "dtype": "number",
          "std": 0.09838699100999075,
          "min": 0.0,
          "max": 0.2,
          "num_unique_values": 4,
          "samples": [0.0],
          "semantic_type": ""
        }
      },
      {
        "column": "ExitRates",
        "properties": {
          "dtype": "number",
          "std": 0.06496152707564687,
          "min": 0.05,
          "max": 0.2,
          "num_unique_values": 4,
          "samples": [0.1],
          "semantic_type": ""
        }
      },
      {
        "column": "PageValues",
        "properties": {
          "dtype": "number",
          "std": 0.0,
          "min": 0.0,
          "max": 0.0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": ""
        }
      }
    ]
  }
}

```

```

    "num_unique_values": 1, "samples": [0.0
], "semantic_type": "", "description": """",
  }, {
    "column": "SpecialDay", "properties": {
      "dtype": "number", "std": 0.0,
      "min": 0.0, "max": 0.0
    }
  },
  "num_unique_values": 1, "samples": [0.0
], "semantic_type": "", "description": """",
  }, {
    "column": "Month", "properties": {
      "dtype": "category", "num_unique_values": 1,
      "samples": ["Feb"]
    }
  },
  "semantic_type": "", "description": """",
  "column": "OperatingSystems", "properties": {
    "dtype": "number", "std": 1,
    "min": 1, "max": 4
  }
}, "num_unique_values": 4, "samples": [2
], "semantic_type": "", "description": """",
  }, {
    "column": "Browser", "properties": {
      "dtype": "number", "std": 0,
      "min": 1, "max": 3
    }
  },
  "num_unique_values": 3, "samples": [1
], "semantic_type": "", "description": """",
  }, {
    "column": "Region", "properties": {
      "dtype": "number", "std": 3,
      "min": 1, "max": 9
    }
  },
  "num_unique_values": 3, "samples": [1
], "semantic_type": "", "description": """",
  }, {
    "column": "TrafficType", "properties": {
      "dtype": "number", "std": 1,
      "min": 1, "max": 4
    }
  },
  "num_unique_values": 4, "samples": [2
], "semantic_type": "", "description": """",
  }, {
    "column": "VisitorType", "properties": {
      "dtype": "category", "num_unique_values": 1,
      "samples": ["Returning_Visitor"]
    }
  },
  "semantic_type": "", "description": """",
  "column": "Weekend", "properties": {
    "dtype": "boolean", "num_unique_values": 2,
    "samples": [true]
  }
}, "semantic_type": "", "description": """",
  }, {
    "column": "Revenue", "properties": {
      "dtype": "boolean", "num_unique_values": 1,
      "samples": [false]
    }
  },
  "semantic_type": "", "description": ""}
], "type": "dataframe"

```

2. Esquema de columnas y saneamiento de tipos

```
# ----- 2) Columnas (según tu esquema en inglés) -----
num_cols = [

```

```

"Administrative", "Administrative_Duration", "Informational", "Informational_Duration",
"ProductRelated", "ProductRelated_Duration", "BounceRates", "ExitRates", "PageValues", "SpecialDay"
]
cat_cols = [
"Month", "OperatingSystems", "Browser", "Region", "TrafficType", "VisitorType", "Weekend"
]
target = "Revenue"

# Seguridad: coerción de tipos, mapeo booleanos 'TRUE/FALSE'
bool_map = {"TRUE": True, "FALSE": False, "True": True, "False": False,
            "true": True, "false": False}
if "Weekend" in df.columns and df["Weekend"].dtype == object:
    df["Weekend"] = df["Weekend"].map(bool_map).fillna(df["Weekend"])
if "Revenue" in df.columns and df["Revenue"].dtype == object:
    df["Revenue"] = df["Revenue"].map(bool_map).fillna(df["Revenue"])

# Coerción numérica segura
for c in num_cols:
    if c in df.columns:
        df[c] = pd.to_numeric(df[c], errors="coerce")

# Seleccionar X, y
use_cols = [c for c in num_cols + cat_cols if c in df.columns]
missing = [c for c in num_cols + cat_cols + [target] if c not in df.columns]
if missing:
    print("⚠ Faltan columnas en el CSV:", missing)

X = df[use_cols].copy()
y = df[target].astype(int) # Revenue -> 0/1

# Rellenar NaN en numéricas con mediana
for c in [c for c in num_cols if c in X.columns]:
    X[c] = X[c].fillna(X[c].median())

```

3. Partición estratificada de datos

```

# ----- 3) Split estratificado -----
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

```

4. Preprocesamiento y definición de modelos

```
# ----- 4) Preprocesamiento + Modelos -----
numeric_features = [c for c in num_cols if c in X.columns]
categorical_features = [c for c in cat_cols if c in X.columns]

preprocess = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numeric_features),
        ("cat", OneHotEncoder(handle_unknown="ignore"),
categorical_features),
    ],
    remainder="drop",
)

# Modelo 1: Regresión Logística
log_reg = Pipeline(steps=[
    ("prep", preprocess),
    ("clf", LogisticRegression(max_iter=1000,
class_weight="balanced"))
])

# Modelo 2: Random Forest (baseline árboles)
rf = Pipeline(steps=[
    ("prep", preprocess),
    ("clf", RandomForestClassifier(
        n_estimators=300, max_depth=None, random_state=42,
class_weight="balanced"
    )))
])

models = {
    "LogisticRegression": log_reg,
    "RandomForest": rf
}
```

5. Función de evaluación y visualización

```
# ----- 5) Entrenamiento + Métricas en Test -----
def evaluate(model, X_train, y_train, X_test, y_test, name="model"):
    model.fit(X_train, y_train)
    proba = model.predict_proba(X_test)[:, 1]
    preds_default = (proba >= 0.5).astype(int)

    acc = accuracy_score(y_test, preds_default)
    f1 = f1_score(y_test, preds_default)
    roc = roc_auc_score(y_test, proba)
    ap = average_precision_score(y_test, proba) # PR AUC

    print(f"\n== {name} ==")
```

```

print(f"Accuracy : {acc:.4f}")
print(f"F1       : {f1:.4f}")
print(f"ROC AUC   : {roc:.4f}")
print(f"PR AUC    : {ap:.4f}")
print("\nClassification report (umbral 0.5):")
print(classification_report(y_test, preds_default, digits=4))

# Matriz de confusión
cm = confusion_matrix(y_test, preds_default)
print("Confusion matrix:\n", cm)

# Curva ROC
fpr, tpr, _ = roc_curve(y_test, proba)
plt.figure(figsize=(5,4))
plt.plot(fpr, tpr, label=f"ROC AUC={roc:.3f}")
plt.plot([0,1],[0,1], '--', alpha=0.7)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title(f"ROC - {name}")
plt.legend()
plt.tight_layout()
plt.show()

# Curva Precision-Recall
prec, rec, thr = precision_recall_curve(y_test, proba)
plt.figure(figsize=(5,4))
plt.plot(rec, prec, label=f"PR AUC={ap:.3f}")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title(f"Precision-Recall - {name}")
plt.legend()
plt.tight_layout()
plt.show()

# ---- Búsqueda de umbral: maximizar F1 ----
best_f1, best_t = 0, 0.5
for t in np.linspace(0.05, 0.95, 19):
    p = (proba >= t).astype(int)
    f1_t = f1_score(y_test, p)
    if f1_t > best_f1:
        best_f1, best_t = f1_t, t
print(f"Mejor umbral por F1: t={best_t:.2f}, F1={best_f1:.4f}")

return {"acc": acc, "f1": f1, "roc_auc": roc, "pr_auc": ap,
"best_thr": best_t, "best_f1": best_f1}

```

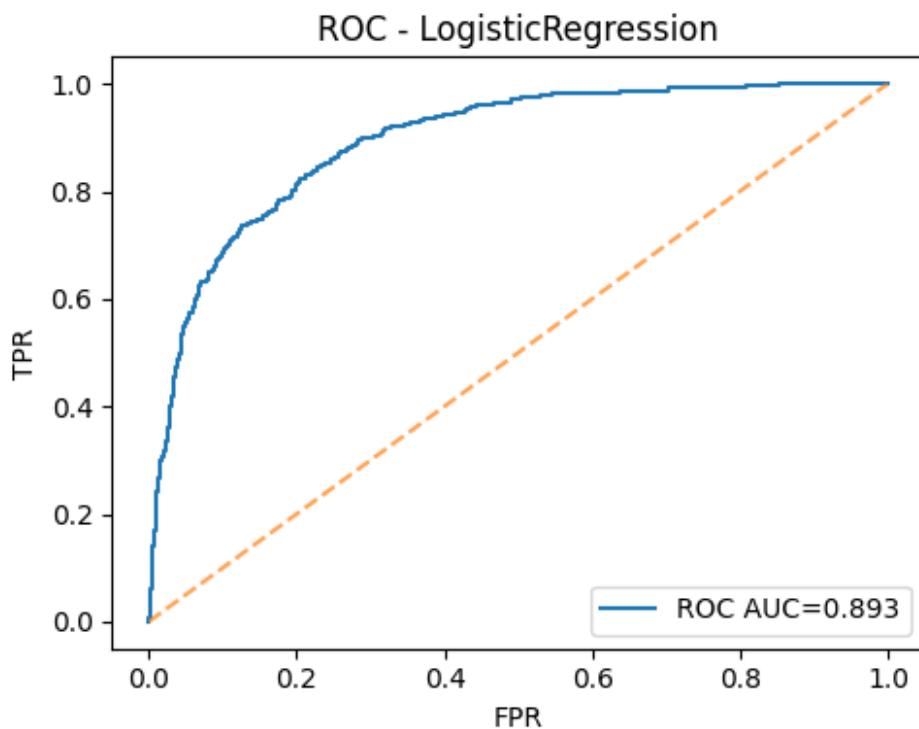
6. Entrenamiento y reporte por modelo

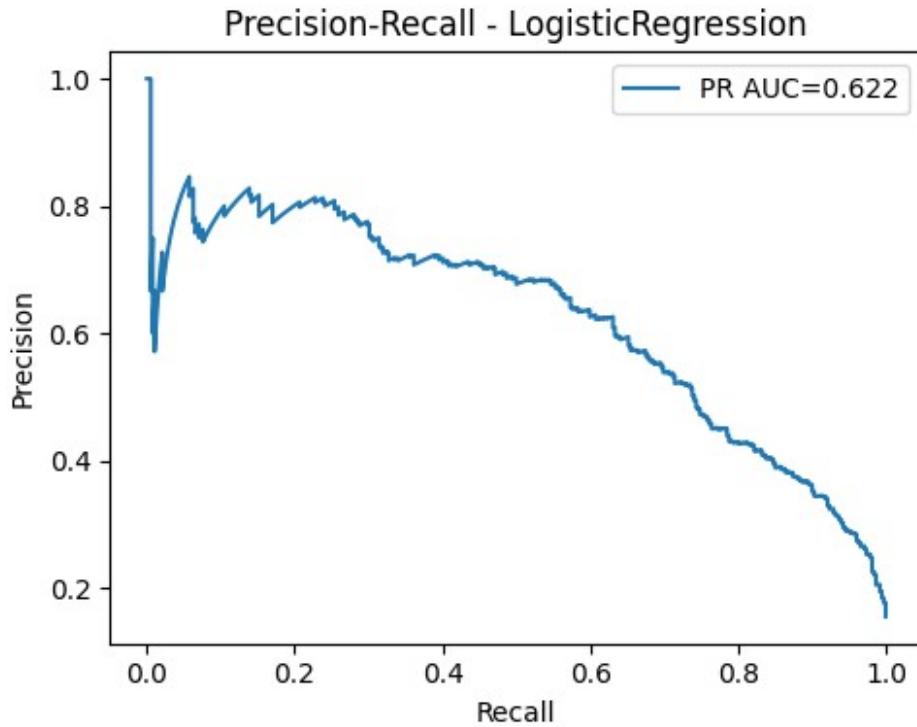
```
results = {}
for name, mdl in models.items():
    results[name] = evaluate(mdl, X_train, y_train, X_test, y_test,
name=name)

==== LogisticRegression ====
Accuracy : 0.8410
F1        : 0.5917
ROC AUC   : 0.8932
PR AUC    : 0.6224

Classification report (umbral 0.5):
precision    recall   f1-score   support
      0       0.9481    0.8589    0.9013    2084
      1       0.4913    0.7435    0.5917     382
      accuracy           0.8410    2466
      macro avg       0.7197    0.8012    0.7465    2466
      weighted avg     0.8773    0.8410    0.8533    2466

Confusion matrix:
[[1790  294]
 [ 98  284]]
```





Mejor umbral por F1: t=0.60, F1=0.6183

```
==== RandomForest ====
Accuracy : 0.8974
F1        : 0.5965
ROC AUC   : 0.9208
PR AUC    : 0.7234
```

```
Classification report (umbral 0.5):
precision    recall    f1-score   support
      0       0.9122    0.9722    0.9412     2084
      1       0.7633    0.4895    0.5965      382

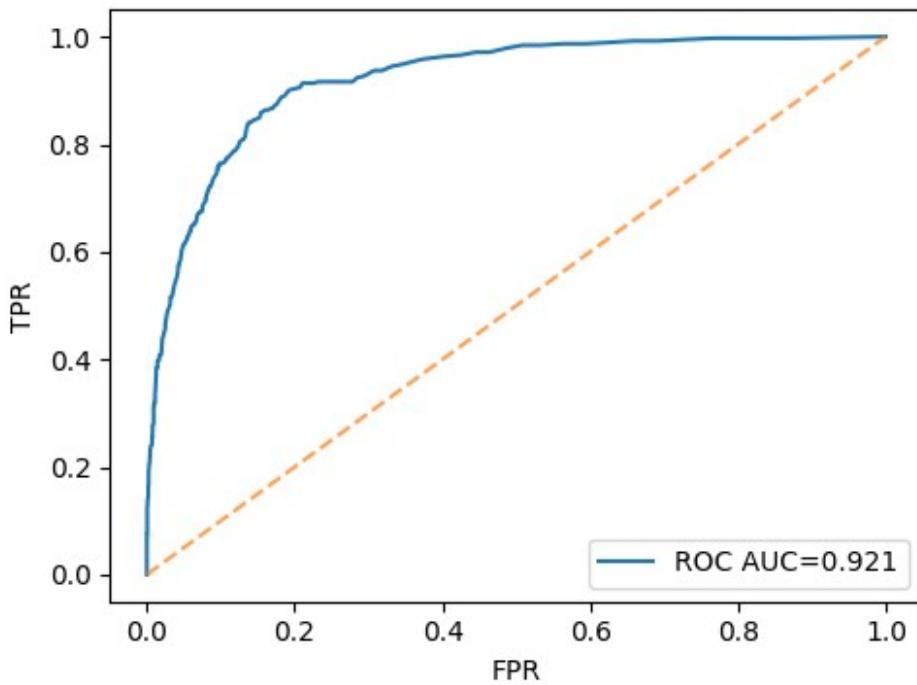
```

	precision	recall	f1-score	support
0	0.9122	0.9722	0.9412	2084
1	0.7633	0.4895	0.5965	382
accuracy			0.8974	2466
macro avg	0.8377	0.7308	0.7689	2466
weighted avg	0.8891	0.8974	0.8878	2466

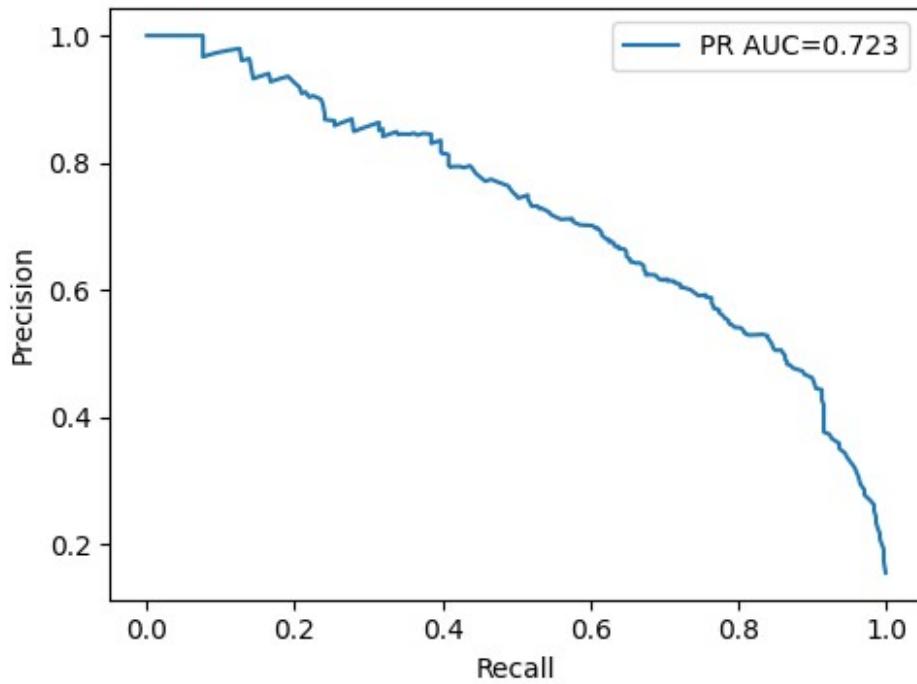
Confusion matrix:

```
[[2026  58]
 [ 195 187]]
```

ROC - RandomForest



Precision-Recall - RandomForest



Mejor umbral por F1: t=0.30, F1=0.6586

7. Validación cruzada (ROC AUC CV5)

```
# ----- 6) Validación cruzada (ROC AUC) -----
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
for name, mdl in models.items():
    scores = cross_val_score(mdl, X, y, cv=cv, scoring="roc_auc")
    print(f"\n{name} - ROC AUC CV5: {scores.mean():.4f} ± {scores.std():.4f}")

LogisticRegression - ROC AUC CV5: 0.9030 ± 0.0077
RandomForest - ROC AUC CV5: 0.9260 ± 0.0048
```

8. Interpretabilidad: coeficientes de la logística

```
# ----- 7) (Opcional) Ver coeficientes de la logística -----
# Para leer coeficientes de la LR, ajustamos y extraemos columnas
# transformadas:
log_reg.fit(X_train, y_train)

# Obtener nombres de columnas después del preprocesamiento
ohe = log_reg.named_steps["prep"].named_transformers_["cat"]
num_names = numeric_features
cat_names = list(ohe.get_feature_names_out(categorical_features)) if
categorical_features else []
feat_names = num_names + cat_names

clf = log_reg.named_steps["clf"]
coefs = pd.Series(clf.coef_.ravel(),
index=feat_names).sort_values(key=np.abs, ascending=False)
print("\nTop coeficientes (|peso|) - LogisticRegression:")
display(coefs.head(20).to_frame("coef"))

Top coeficientes (|peso|) - LogisticRegression:

{"summary": "\n  \"name\": \"display(coefs\", \"rows\": 20,\n  \"fields\": [\n    {\n      \"column\": \"coef\", \"\n      \"properties\": {\n        \"dtype\": \"number\", \"\n        \"std\": 0.9594114793611253, \"\n        \"min\": -1.2732857577802994, \"\n        \"max\": 2.3074413045461784, \"\n        \"num_unique_values\": 20, \"\n        \"samples\": [\n          2.3074413045461784, \"\n          0.424640468403154, \"\n          0.4919118142891404\n        ], \"\n        \"semantic_type\": \"\", \"\n        \"description\": \"\"\n      }\n    }\n  ]\n}, \"type\": \"dataframe\"}
```