

# Modelo de regresión

Integrante: Luis Keny Lucero Balvin

```
# Archivo: clientes_intereses_1000_targets.csv
# X: primeras 12 columnas (intereses 1-10)
# Y1: convirtio_30d (CLASIFICACIÓN) | Y2: ordenes_90d (REGRESIÓN DE
# CONTEO)

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression, PoissonRegressor
from sklearn.metrics import (
    accuracy_score, f1_score, roc_auc_score, classification_report,
    confusion_matrix
)
```

## 1. Dependencias para VIF

```
# Para VIF
!pip -q install statsmodels
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import
variance_inflation_factor
```

## 2. Parámetro de entrada (ruta del CSV)

```
CSV = "clientes_intereses_1000_targets.csv"
```

## 3. Carga del dataset

```
# 1) Carga
if not os.path.exists(CSV):
    from google.colab import files
    print(f"No encontré '{CSV}'. Súbelo ahora...")
    uploaded = files.upload()
    if CSV not in uploaded:
        CSV = list(uploaded.keys())[0]
        print("Usando:", CSV)

df = pd.read_csv(CSV)
```

## 4. Esquema de columnas esperadas

```
# 2) Columnas esperadas
feature_cols = [
    "importancia_garantia", "sensibilidad_precio", "rapidez_envio", "calidad_producto",
    "devolucion_flexible", "confianza_reviews", "sostenibilidad", "marca_importancia",
    "variedad_metodos_pago", "personalizacion", "soporte_postventa", "fidelidad_programa"
]
y_cls = "convirtio_30d"
y_cnt = "ordenes_90d"

missing = [c for c in feature_cols+[y_cls, y_cnt] if c not in df.columns]
if missing:
    raise ValueError(f"Faltan columnas en el CSV: {missing}")
```

## 5. Limpieza mínima y coerción de tipos

```
# 3) Limpieza rápida
X = df[feature_cols].copy()
y1 = df[y_cls].copy()
y2 = df[y_cnt].copy()

# Forzar numérico y tratar NaN por si acaso
for c in feature_cols:
    X[c] = pd.to_numeric(X[c], errors="coerce")
y1 = pd.to_numeric(y1, errors="coerce")
y2 = pd.to_numeric(y2, errors="coerce")
```

## 6. Diagnóstico de faltantes y relleno

```
# Reporte de faltantes
null_report = pd.DataFrame({
    "faltantes_X": X.isna().sum(),
})
null_report.loc["convirtio_30d", "faltantes_Y1"] = y1.isna().sum()
null_report.loc["ordenes_90d", "faltantes_Y2"] = y2.isna().sum()
print("==== Faltantes ===")
display(null_report.fillna(0).astype(int))

# Relleno conservador con medianas
X = X.fillna(X.median())
```

```

y1 = y1.fillna(y1.median())
y2 = y2.fillna(y2.median())

==== Faltantes ====

{
  "summary": {
    "name": "y2 = y2",
    "rows": 14,
    "fields": [
      {
        "column": "faltantes_X",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "faltantes_Y1",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "faltantes_Y2",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 0,
          "num_unique_values": 1,
          "samples": [0],
          "semantic_type": "",
          "description": ""
        }
      }
    ],
    "type": "dataframe"
  }
}

```

## 7. Columnas constantes o casi constantes

```

# 4) Columnas constantes o con varianza ~0
nuniq = X.nunique()
const_cols = nuniq[nuniq <= 1].index.tolist()
print("Columnas constantes/varianza cero:", const_cols if const_cols else "Ninguna")

Columnas constantes/varianza cero: Ninguna

```

## 8. Chequeo de rangos esperados (1–10)

```

# 5) Rango de features (esperado 1–10)
print("==== Rango observado de features ===")
display(pd.DataFrame({"min": X.min(), "max": X.max()}))

==== Rango observado de features ===

{
  "summary": {
    "name": "display(pd",
    "rows": 12,
    "fields": [
      {
        "column": "min",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 3,
          "num_unique_values": 4,
          "samples": [2, 2, 3, 3],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "max",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 10,
          "num_unique_values": 1,
          "samples": [10],
          "semantic_type": "",
          "description": ""
        }
      }
    ],
    "type": "dataframe"
  }
}

```

```

  "std": 0, "min": 8, "max": 10,
  "num_unique_values": 3, "samples": [9,
  10, 8], "semantic_type": "\",
  "description": """
  }], "type": "dataframe"
}

```

## 9. Multicolinealidad (VIF)

```

# 6) Multicolinealidad (VIF)
# Escalamos y añadimos constante para VIF
Xs = pd.DataFrame(StandardScaler().fit_transform(X),
columns=feature_cols)
Xs_const = sm.add_constant(Xs, has_constant="add")
vif = pd.Series(
    [variance_inflation_factor(Xs_const.values, i) for i in
range(Xs_const.shape[1])],
    index=["const"] + feature_cols,
    name="VIF")
print("\n==== VIF (multicolinealidad) ====")
display(vif.to_frame())

==== VIF (multicolinealidad) ====

{
  "summary": {
    "name": "display(vif\", \"rows\": 13,
    "fields": [
      {
        "column": "VIF",
        "properties": {
          "dtype": "number",
          "std": 0.9984657436771056,
          "min": 0.9999999999999998,
          "max": 4.994573684954479,
          "num_unique_values": 13,
          "samples": [
            3.1149816909519226,
            2.2486339122331422,
            0.9999999999999998
          ],
          "semantic_type": "\",
          "description": """
          }
        ]
      }
    ],
    "type": "dataframe"
}

```

## 10. Exploratorio: histogramas y correlación


```


# 7) Distribuciones y correlación rápida
axes = X.hist(bins=10, figsize=(14,10))
plt.suptitle("Histogramas - Features", y=1.02); plt.tight_layout();
plt.show()

corr = X.corr(method="pearson")
plt.figure(figsize=(9,7))
im = plt.imshow(corr, cmap="RdBu_r", vmin=-1, vmax=1)
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.index)), corr.index)
plt.colorbar(im, fraction=0.046, pad=0.04)
plt.title("Matriz de correlación (Pearson)")
plt.tight_layout(); plt.show()

```

Histogramas - Features





## 11. Naturaleza de los objetivos (Y)

```
# 8) Naturaleza de las Y
print("\n==== Naturaleza de las Y ====")
cls_rate = y1.mean()
zeros_rate = (y2 == 0).mean()
print(f"convertio_30d -> tasa de positivos = {cls_rate:.3f}")
print(f"ordenes_90d    -> prop. de ceros   = {zeros_rate:.3f}, "
     f"media={y2.mean():.3f}, var={y2.var():.3f}")

==== Naturaleza de las Y ====
convertio_30d -> tasa de positivos = 0.450
ordenes_90d    -> prop. de ceros   = 0.455, media=0.966, var=1.578
```

## 12. Señal base (clasificación)

```
# 9) Señal base rápida
# Clasificación: Regresión Logística
X_tr, X_te, y1_tr, y1_te = train_test_split(X, y1, test_size=0.2,
random_state=42, stratify=y1)
scaler = StandardScaler().fit(X_tr)
X_tr_s, X_te_s = scaler.transform(X_tr), scaler.transform(X_te)

lr = LogisticRegression(max_iter=1000, class_weight="balanced")
lr.fit(X_tr_s, y1_tr)
p = lr.predict_proba(X_te_s)[:,1]
pred = (p >= 0.5).astype(int)
```

## 13. Métricas y reporte (clasificación)

```
# === Línea base: Clasificación (Logistic Regression) ===
print("\n== Línea base: Clasificación (Logistic Regression) ==")

# Probabilidades y predicción por defecto (umbral 0.5)
p = lr.predict_proba(X_te_s)[:, 1]
pred = (p >= 0.5).astype(int)

# Métricas (¡ojo! usar formato o round(), no .round() sobre float)
roc = roc_auc_score(y1_te, p)
f1 = f1_score(y1_te, pred)
acc = accuracy_score(y1_te, pred)

print(f"ROC AUC: {roc:.3f}")
print(f"F1 : {f1:.3f}")
print(f"Accuracy: {acc:.3f}")

print("\nClassification report (umbral 0.5):")
print(classification_report(y1_te, pred, digits=3))

cm = confusion_matrix(y1_te, pred)
print("Confusion matrix:\n", cm)

== Línea base: Clasificación (Logistic Regression) ==
ROC AUC: 0.701
F1 : 0.606
Accuracy: 0.655

Classification report (umbral 0.5):
precision    recall   f1-score   support
      0       0.678      0.709      0.693      110
      1       0.624      0.589      0.606       90
```

accuracy			0.655	200
macro avg	0.651	0.649	0.650	200
weighted avg	0.654	0.655	0.654	200

Confusion matrix:

```
[[78 32]
 [37 53]]
```

## 14. Señal base (conteos Poisson)

```
# Conteos: Poisson
X_tr2, X_te2, y2_tr, y2_te = train_test_split(X, y2, test_size=0.2,
random_state=42)
scaler2 = StandardScaler().fit(X_tr2)
X_tr2_s, X_te2_s = scaler2.transform(X_tr2), scaler2.transform(X_te2)

pr = PoissonRegressor(alpha=1e-6, max_iter=1000) # base
pr.fit(X_tr2_s, y2_tr)
y2_pred = pr.predict(X_te2_s)
```

## 15. Métricas de conteo y sugerencia de familia

```
rmse = np.sqrt(np.mean((y2_te - y2_pred)**2))
mae = np.mean(np.abs(y2_te - y2_pred))
print("\n==== Línea base: Regresión de conteo (Poisson) ===")
print("RMSE:", round(rmse,3), " | MAE:", round(mae,3))
print("Media y2_test:", round(y2_te.mean(),3), " | Var y2_test:",
round(y2_te.var(),3))
print("Sugerencia modelo de conteo:",
      "Negativo Binomial/Zero-Inflated" if (y2_te.var() > y2_te.mean()*2 or
zeros_rate > 0.4) else "Poisson parece razonable")

==== Línea base: Regresión de conteo (Poisson) ===
RMSE: 0.965 | MAE: 0.711
Media y2_test: 1.075 | Var y2_test: 1.477
Sugerencia modelo de conteo: Negativo Binomial/Zero-Inflated
```